

Lecture 9: Bayesian DSGE Models

Instructor: Fei Tan



@econdoj



@BusinessSchool101



Saint Louis University

Course: Introduction to Bayesian Statistics

Date: January 31, 2026

Introduction

- DYNARE has played a large role in fitting DSGE models
- However, it cannot handle *high-dimensional* DSGE models or models with
 - Student- t shocks, e.g., Chib & Ramamurthy (2014)
 - stochastic volatility, e.g., Justiniano & Primiceri (2008)
- Chib, Shin, & Tan (2021) provide a user-friendly MATLAB toolbox for such models that contains
 - training sample priors
 - efficient sampling of parameters by the TaRB-MH algorithm of Chib & Ramamurthy (2010)
 - fast computation of the marginal likelihood by the method of Chib (1995) and Chib & Jeliazkov (2001)
 - post-estimation tools, e.g., point and density forecasts

The Road Ahead

1. DSGE Model
2. Prior Distribution
3. Posterior Sampling
4. Marginal Likelihood
5. Prediction

Illustration

- Dynamic IS relation

$$\hat{y}_t = \mathbb{E}_t \hat{y}_{t+1} + \hat{g}_t - \mathbb{E}_t \hat{g}_{t+1} - \frac{1}{\tau} (\hat{R}_t - \mathbb{E}_t \hat{\pi}_{t+1} - \mathbb{E}_t \hat{z}_{t+1})$$

- New Keynesian Phillips curve

$$\hat{\pi}_t = \beta \mathbb{E}_t \hat{\pi}_{t+1} + \kappa (\hat{y}_t - \hat{g}_t)$$

- Monetary policy rule

$$\hat{R}_t = \rho_R \hat{R}_{t-1} + (1 - \rho_R) \psi_1 \hat{\pi}_t + (1 - \rho_R) \psi_2 (\hat{y}_t - \hat{g}_t) + \epsilon_{R,t}$$

- Exogenous shocks

$$\hat{g}_t = \rho_G \hat{g}_{t-1} + \epsilon_{G,t}, \quad \hat{z}_t = \rho_Z \hat{z}_{t-1} + \epsilon_{Z,t}$$

- t -innovation with SV: for $s \in \{R, G, Z\}$

$$\epsilon_{s,t} \sim t_\nu(0, e^{h_{s,t}}), \quad h_{s,t} = (1 - \phi_s) \mu_s + \phi_s h_{s,t-1} + \eta_{s,t}, \quad \eta_{s,t} \sim N(0, \omega_s^2)$$

State Space Form

- We rely on Sims' (2001) method to solve for the solution of the structural model, for each value of θ
- After (log) linearizing around the steady state, the structural model is expressed in canonical form as

$$\Gamma_0(\theta)x_t = \Gamma_1(\theta)x_{t-1} + \Psi\epsilon_t + \Pi\eta_t$$

- Applying Sims' method, the unique, bounded solution takes the form

$$x_t = G(\theta)x_{t-1} + M(\theta)\epsilon_t$$

where $G(\theta)$ and $M(\theta)$ are non-linear unspecified functions of θ determined by the solve step

State Space Form (Cont'd)

- Model completed by the measurement equations

$$\begin{pmatrix} \text{YGR}_t \\ \text{INF}_t \\ \text{INT}_t \end{pmatrix} = \begin{pmatrix} \gamma^{(Q)} \\ \pi^{(A)} \\ \pi^{(A)} + r^{(A)} + 4\gamma^{(Q)} \end{pmatrix} + \begin{pmatrix} \hat{y}_t - \hat{y}_{t-1} + \hat{z}_t \\ 4\hat{\pi}_t \\ 4\hat{R}_t \end{pmatrix}$$

- We show results for an extended version of this model that has 51 parameters, 21 variables, 8 shocks, 8 observables, and 1,494 non-Gaussian and nonlinear latent variables

Example

- Set MATLAB directory to the 'DSGE-SVt' folder
- Specify the model, data, and save directories

```
%% Housekeeping
clear
close all
clc

%% User search path & mex files
modpath = ['user' filesep 'ltw17'];
datpath = ['user' filesep 'ltw17' filesep 'data.txt'];
savepath = ['user' filesep 'ltw17'];
spec = tarb_spec([], 'modpath', modpath, 'datpath', datpath, 'savepath', savepath);

OneFileToMexThemAll
```

Training Sample Prior

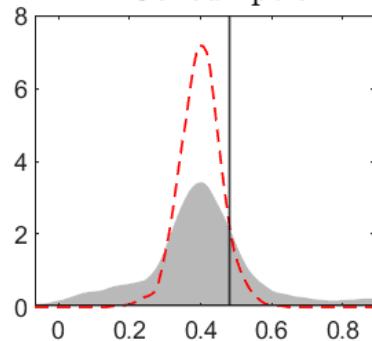
- In high dimensions, formulating an appropriate prior is difficult due to the complex mapping from θ to $G(\theta)$ and $M(\theta)$
- Standard choices often produce prior-sample conflict
- We supply two ways of dealing with this: training sample priors and Student- t family of distributions for location-type parameters
- A sampling the prior function is available to calculate the implied distribution of the outcomes

```
%% Sampling the prior
npd = 10000;           % number of prior draws
T = 200;                % number of periods
sdof = 5;                 % shock degrees of freedom
SamplePrior(npd,T,sdof,savepath)
```

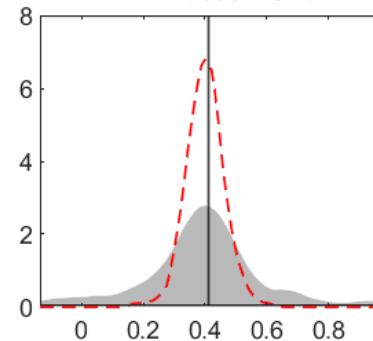
Prior Predictive Distribution

I. Sample mean

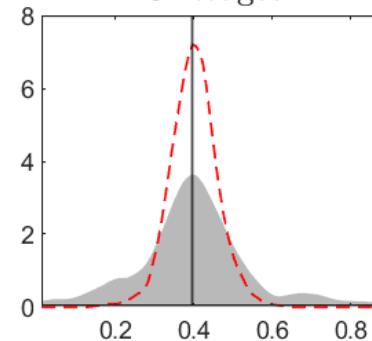
A. Consumption



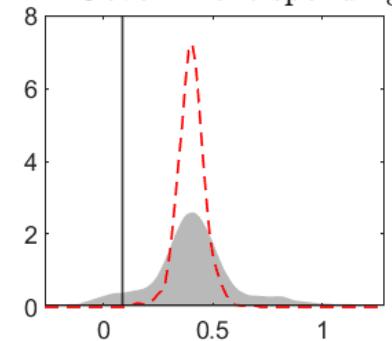
B. Investment



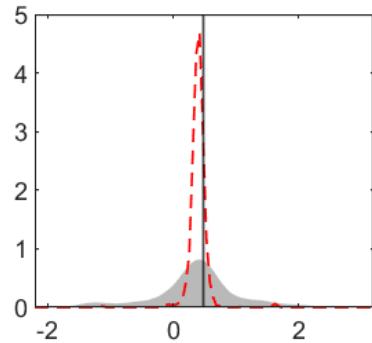
C. Wages



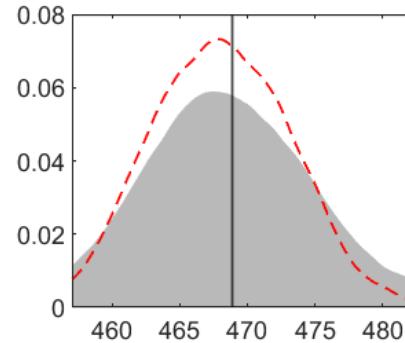
D. Government spending



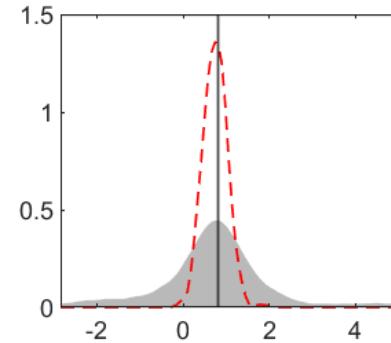
E. Debt



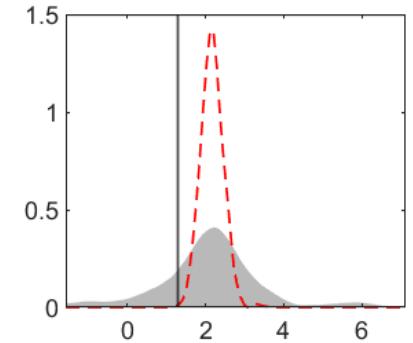
F. Hours worked



G. Inflation



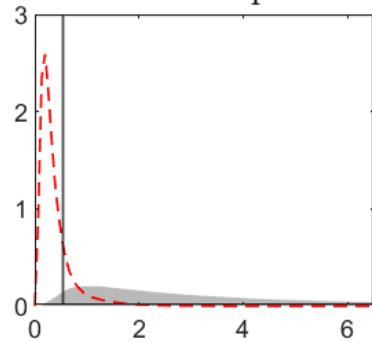
H. Interest rate



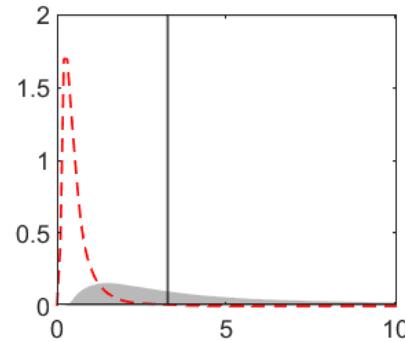
Prior Predictive Distribution (Cont'd)

II. Sample standard deviation

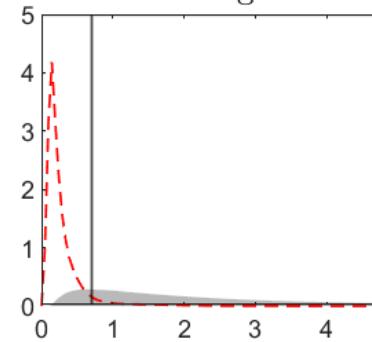
A. Consumption



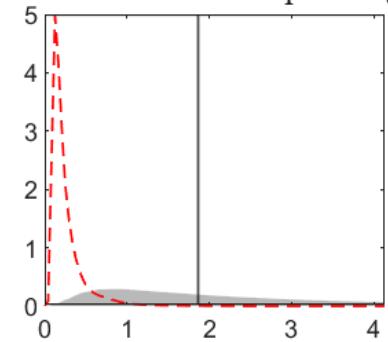
B. Investment



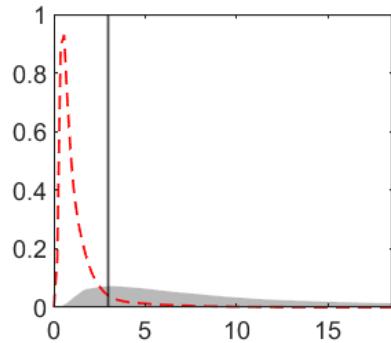
C. Wages



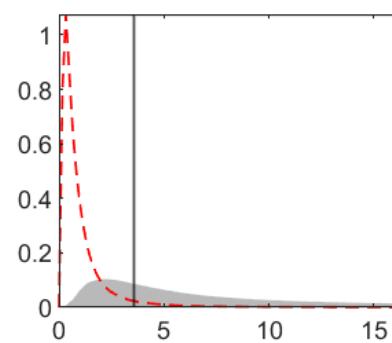
D. Government spending



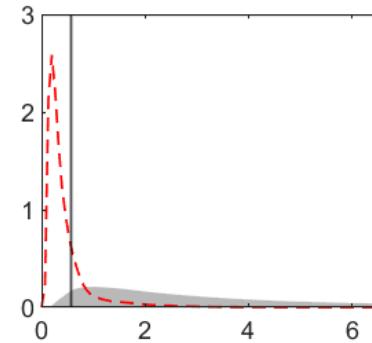
E. Debt



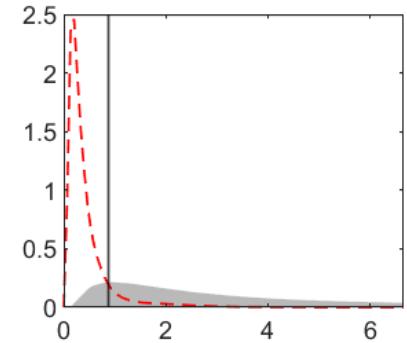
F. Hours worked



G. Inflation



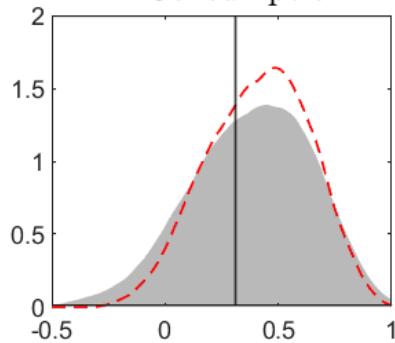
H. Interest rate



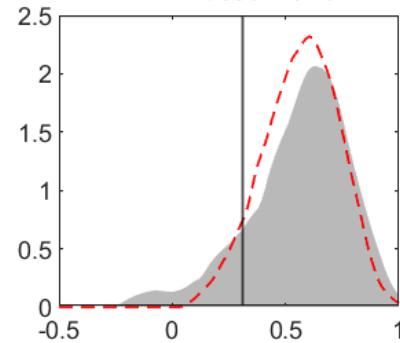
Prior Predictive Distribution (Cont'd)

III. Sample first-order autocorrelation

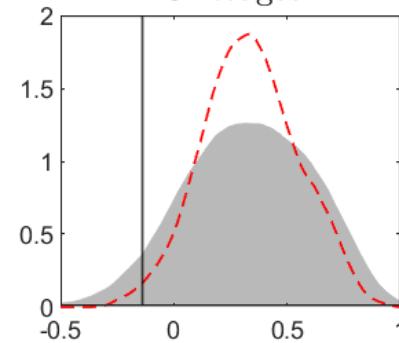
A. Consumption



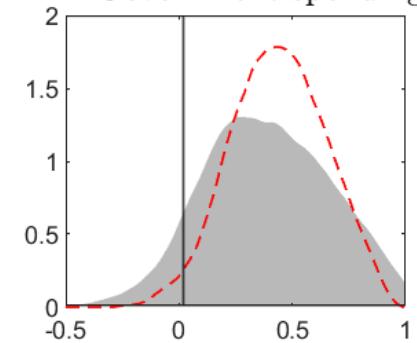
B. Investment



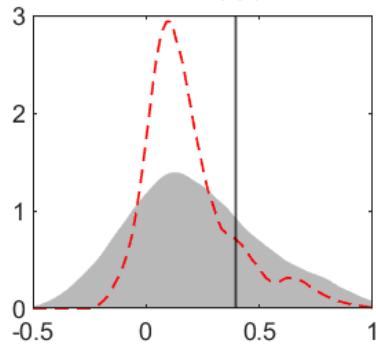
C. Wages



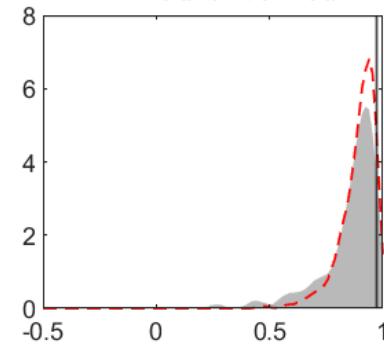
D. Government spending



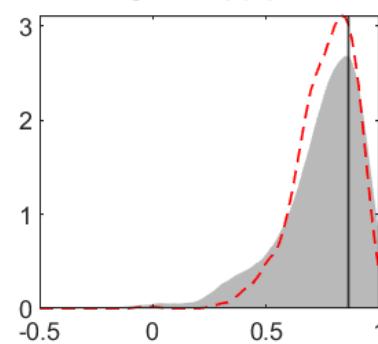
E. Debt



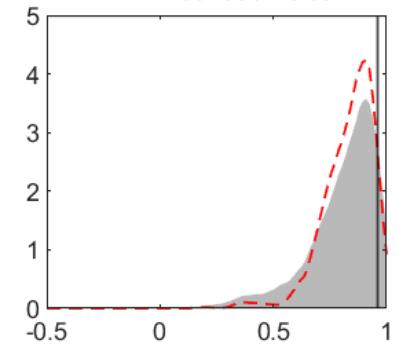
F. Hours worked



G. Inflation



H. Interest rate



Estimation

- The centerpiece of the estimation procedure is the TaRB-MH algorithm of Chib & Ramamurthy (2010)
- Used to draw samples of $(\theta, z_{1:T})$ from

$$\pi(\theta, z_{1:T} | y_{1:T}) \propto f(y_{1:T}, z_{1:T} | \theta) \cdot \pi(\theta) \cdot \mathbf{1}\{\theta \in \Theta_D\}$$

- TaRB-MH is coded up in DYNARE, but the implementation there is somehow not efficient

TaRB-MH MCMC

- Two hallmarks of TaRB-MH
 - randomize the number and components of blocks
 - tailor the proposal density to the posterior location and curvature, i.e.,
$$q_b(\theta_b | y_{1:T}, z_{1:T}, \theta_{-b}) = t_v(\theta_b | \hat{\theta}_b, \hat{V}_b), b = 1, \dots, B,$$
using simulated annealing and Chris Sims' `csminwel`
- The toolbox provides a fast implementation of these steps by
 - randomly tailoring proposal every few iterations
 - MEX loop-intensive functions from C/C++ code
 - `csminwel` to get \hat{V}_b as output
- Although intensive it all happens seamlessly

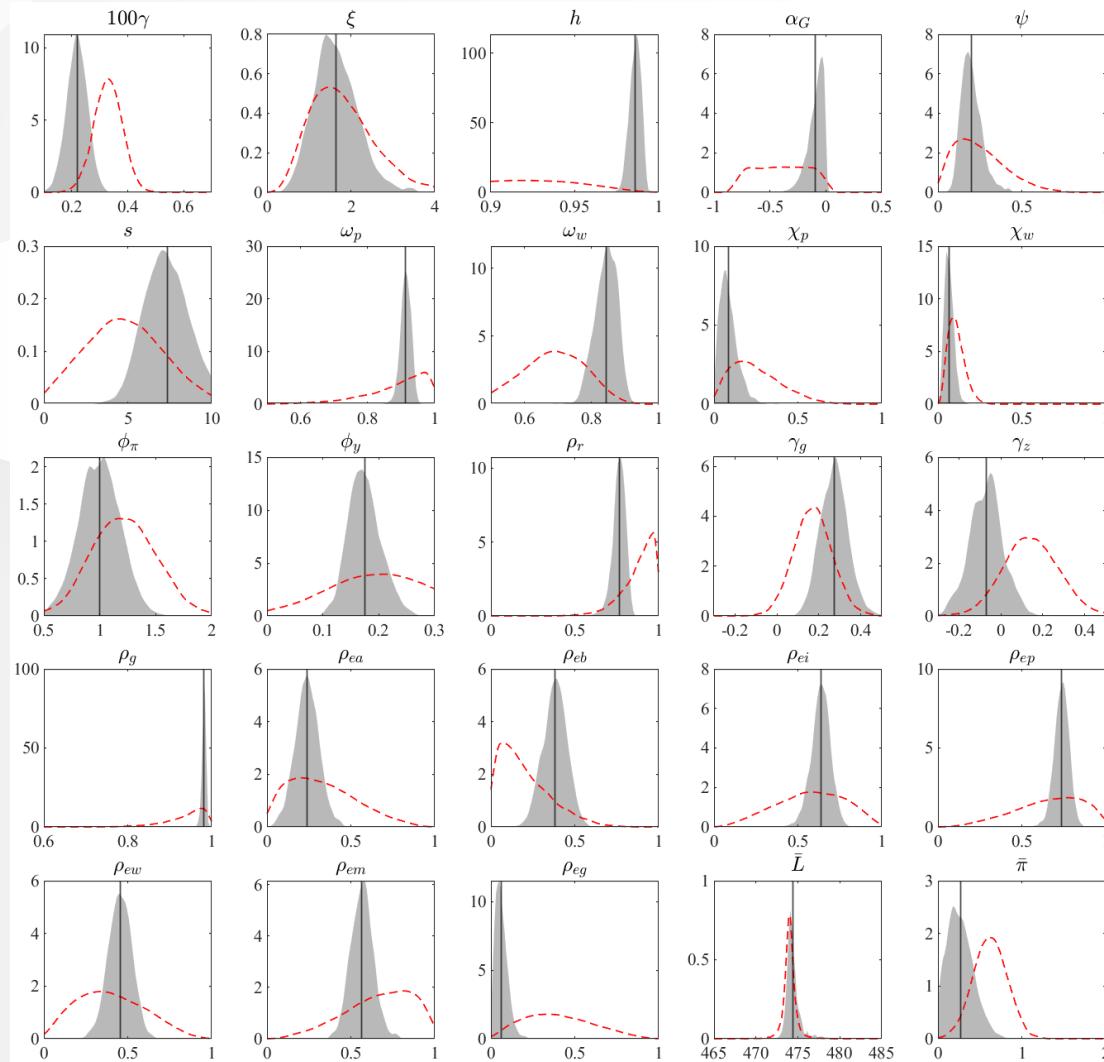
Example

- Set up the TaRB-MH algorithm
- The estimation results are stored in the MATLAB data file `tarb_full.mat`

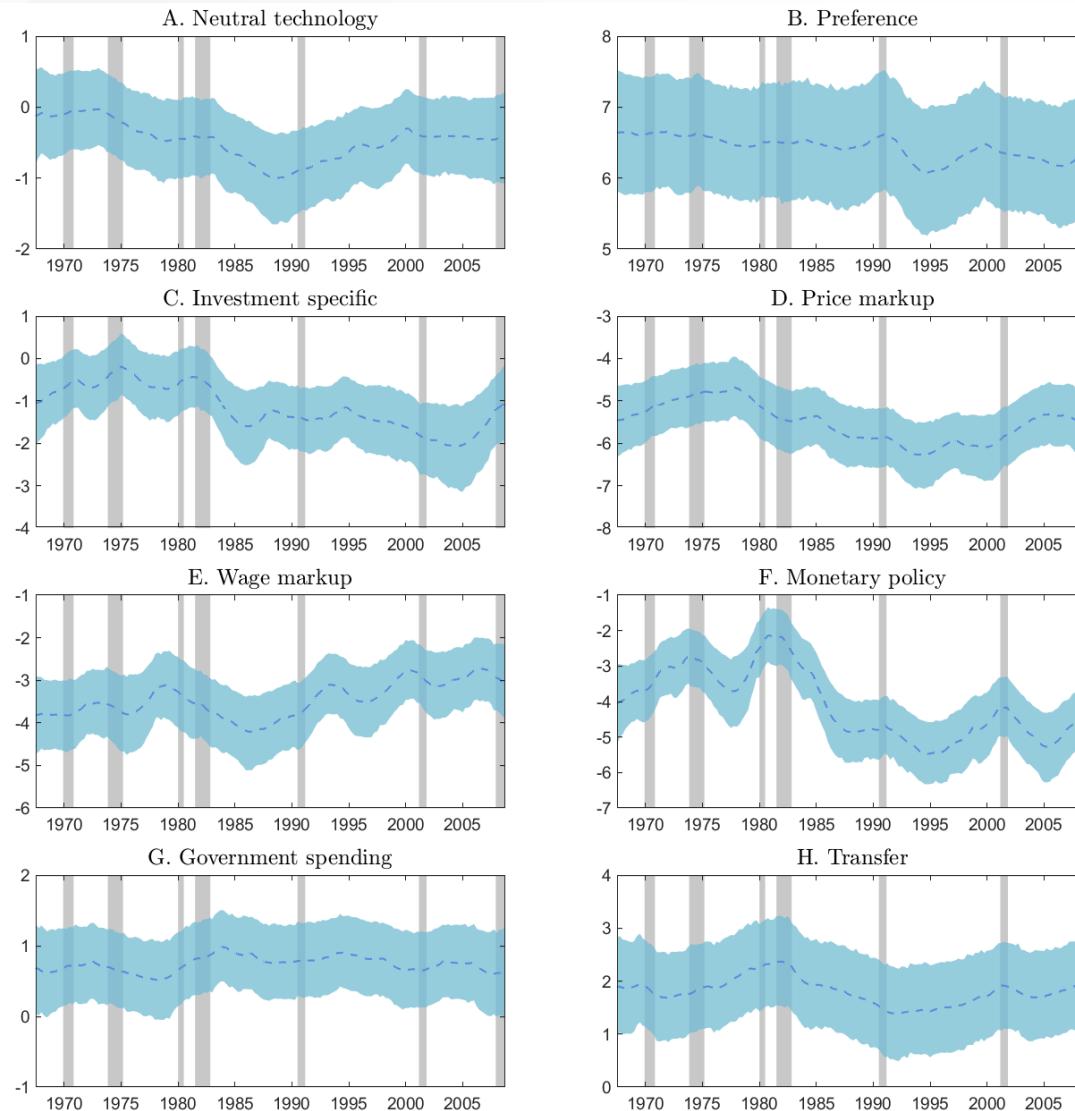
```
%% TaRB-MH (full run)
p = 0.7; % blocking probability
w = 0.5; % tailoring frequency
spec = tarb_spec(spec,'prob',p,'freq',w);

M = 11000; % number of draws
burn = 1000; % number of burn-in
tarb_full(M,burn,spec)
```

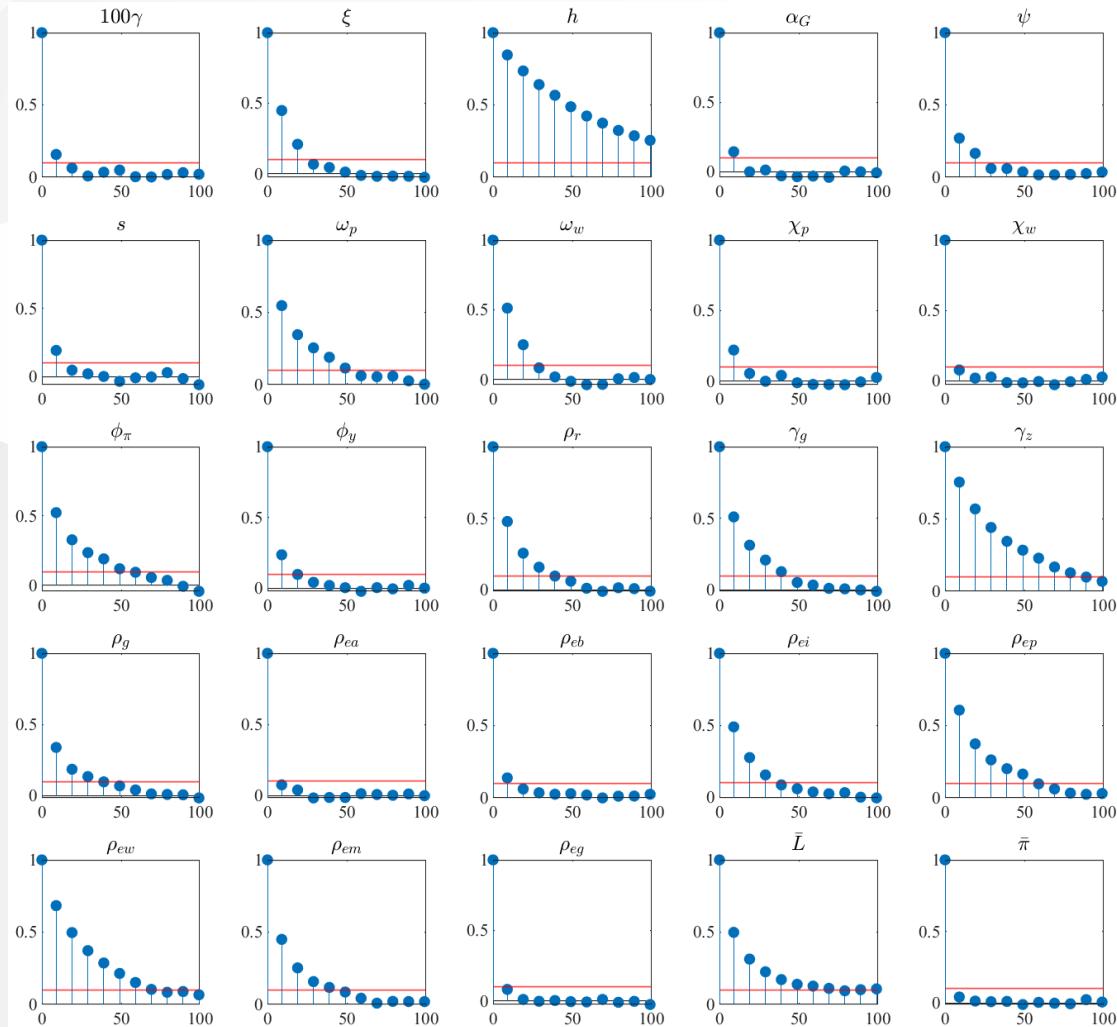
Prior-Posterior



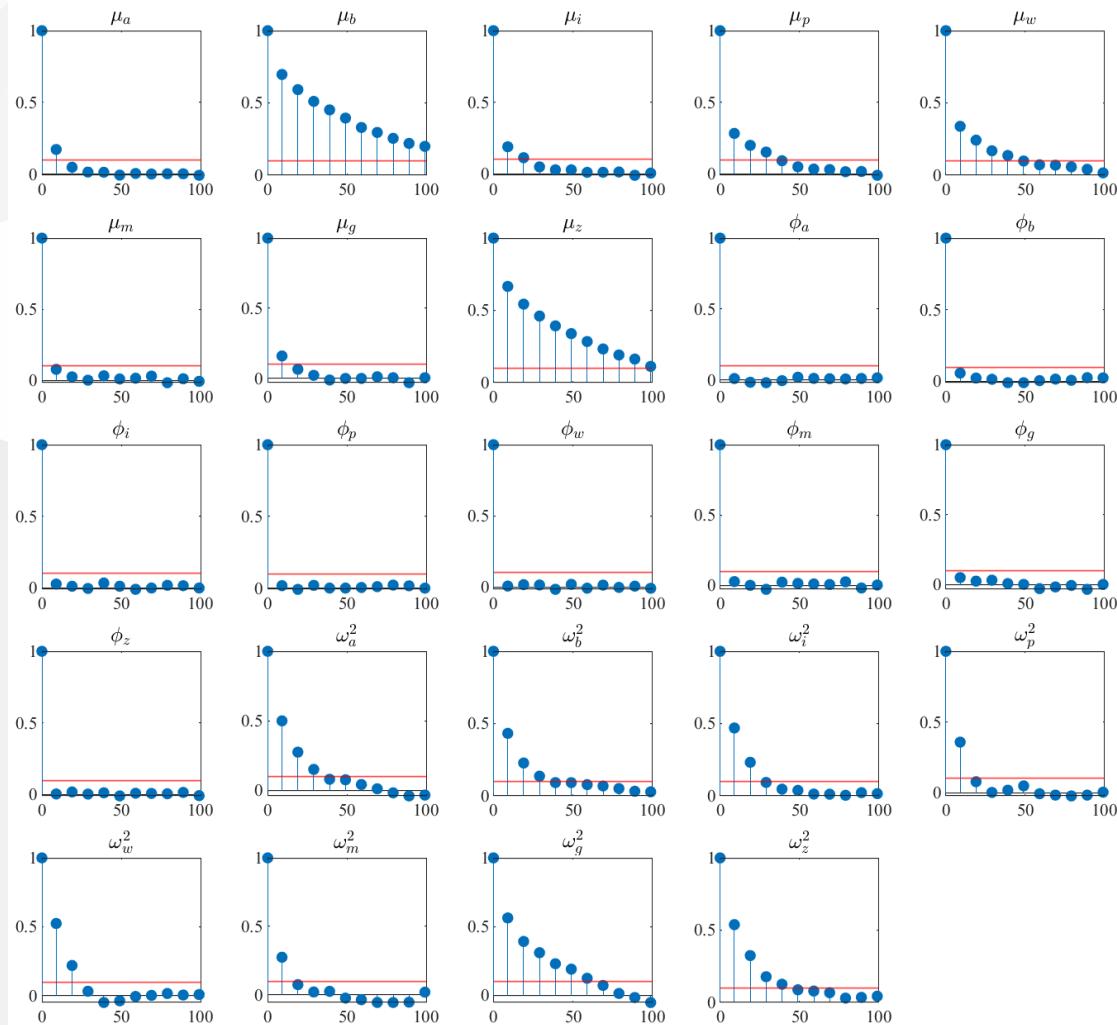
Stochastic Volatility



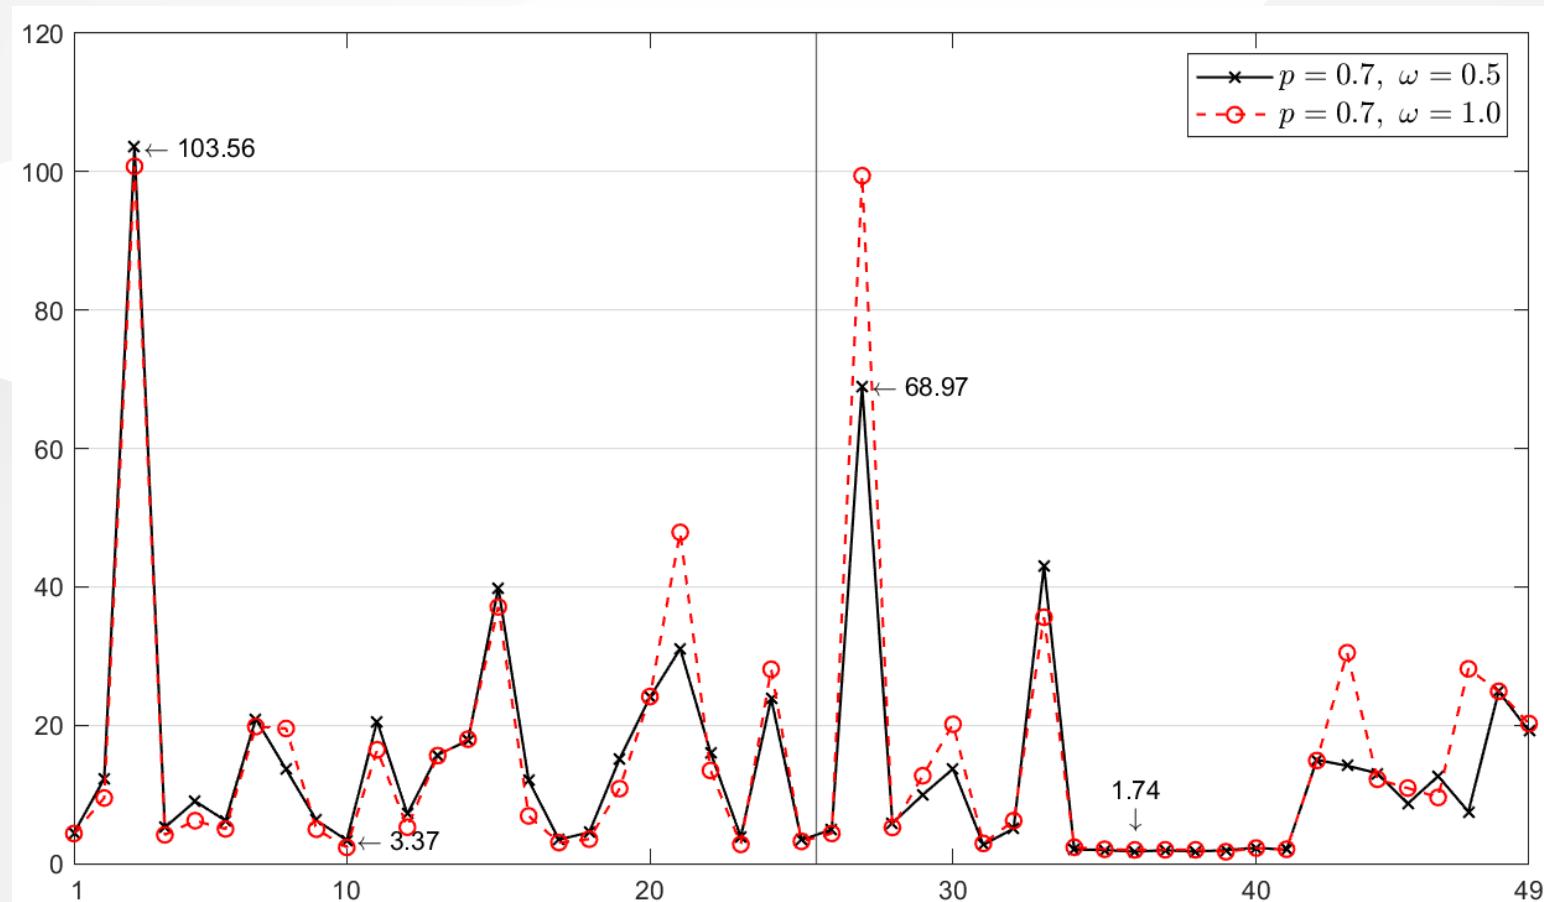
Autocorrelation



Autocorrelation (Cont'd)



Inefficiency Factor



Model Comparison

- Marginal likelihood is computed by a fast implementation of the Chib & Jeliazkov (2001) estimator
- This estimator is based on the identity of Chib (1995)

$$m(y_{1:T}) = \frac{f(y_{1:T}|\theta^*)\pi(\theta^*)}{\pi(\theta^*|y_{1:T})}, \quad \forall \theta^*$$

Model Comparison (Cont'd)

- With multiple block sampling, the posterior ordinate is estimated from the decomposition

$$\pi(\theta^* | y_{1:T}) = \pi(\theta_1^* | y_{1:T})\pi(\theta_2^* | y_{1:T}, \theta_1^*) \cdots \pi(\theta_B^* | y_{1:T}, \theta_1^*, \dots, \theta_{B-1}^*)$$

using several reduced TaRB-MH runs

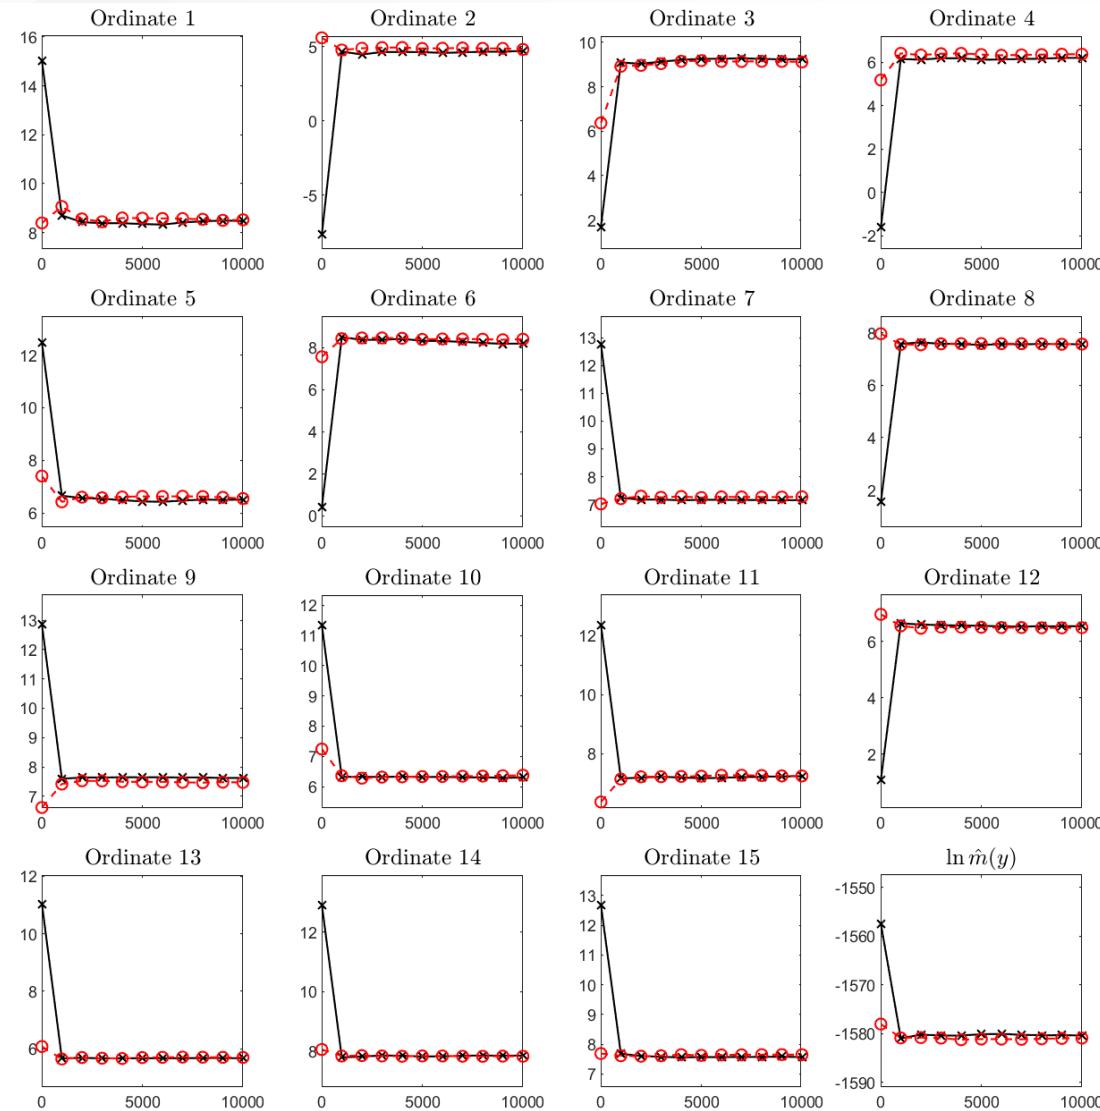
- Key point: all reduced runs can be parallelized for the cost of only one reduced run, regardless of the number of blocks
- This speeds up the calculation enormously

Example

- Set up the TaRB-MH reduced run
- The estimation results are stored in the MATLAB data file `tarb_reduce.mat`

```
%% TaRB-MH (reduced run)
M = 11000; % number of draws
burn = 1000; % number of burn-in
B = 7; % number of blocks
tarb_reduce(M,burn,B,spec)
```

Marginal Likelihood



Simulation Evidence

- Chib-Jeliazkov estimator

DGP 1: regime-M with $v=15$			DGP 2: regime-F with $\phi=0.5$		
v	M	F	ϕ	M	F
30 (light)	4	0	0.1 (weak)	0	9
15 (fat)	15	0	0.5 (moderate)	0	10
5 (heavy)	1	0	0.9 (strong)	0	1

- Modified harmonic mean estimator

DGP 1: regime-M with $v=15$			DGP 2: regime-F with $\phi=0.5$		
v	M	F	ϕ	M	F
30 (light)	0	0	0.1 (weak)	0	0
15 (fat)	0	0	0.5 (moderate)	0	0
5 (heavy)	20	0	0.9 (strong)	0	20

Notes: Number of picks for each model specification.

Post-Estimation Tools

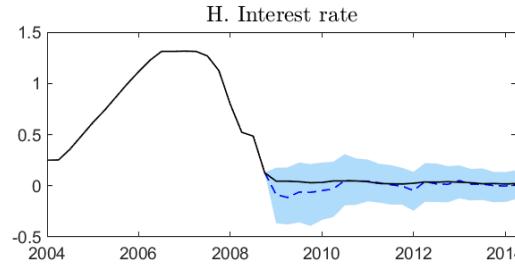
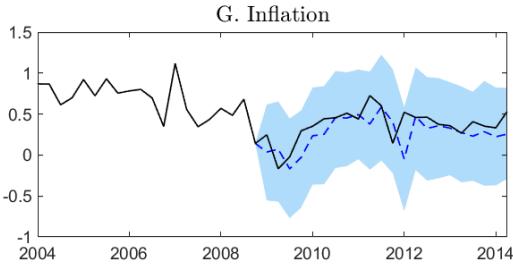
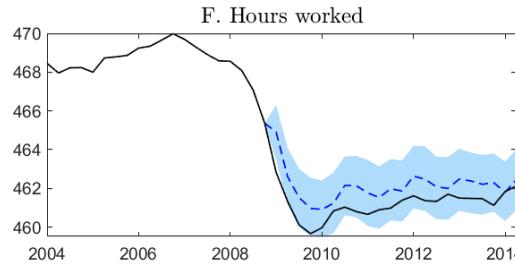
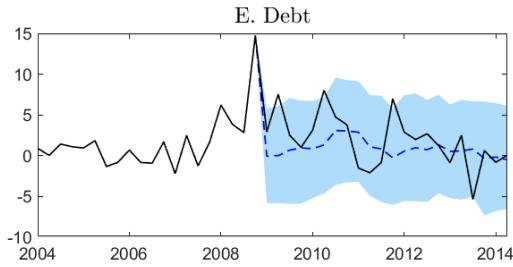
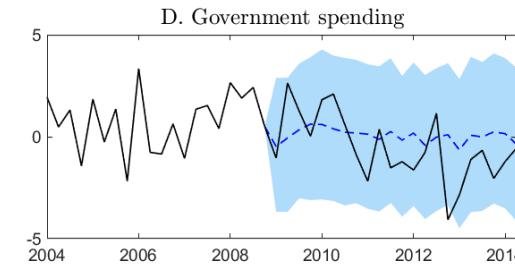
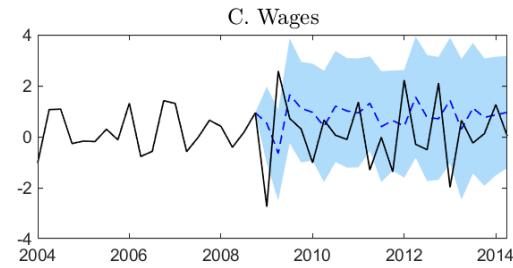
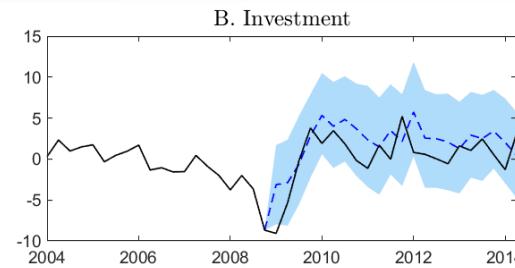
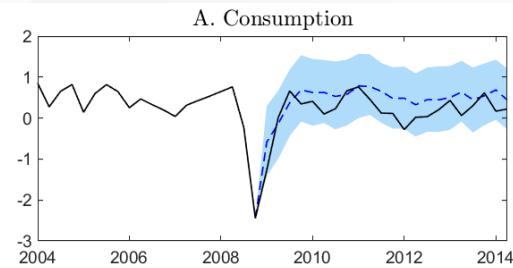
- Toolbox includes post-estimation tools, e.g., functions for conducting impulse response, variance decomposition, and prediction
- For instance, suppose we want the one-quarter-ahead prediction density

$$p(y_{T+1}|y_{1:T}) = \int p(y_{T+1}|y_{1:T}, \theta, z_{T+1}) \cdot \pi(\theta, z_{T+1}|y_{1:T}) d(\theta, z_{1:T})$$

- This is available through the TaRB-MH specification

```
%% Out-of-sample forecast
head = 50;          % training sample size
tail = 22;          % forecasting sample size
h = 1;              % forecasting horizon
spec = tarb_spec([], 'datrange', [head tail], 'pred', h);
```

Out-of-Sample Forecast



Outroduction

- Going beyond DYNARE for Next-Gen DSGE models
- Efficient and fast estimation via TaRB-MH and parallel computing
- MATLAB toolbox: publicly available at github.com/econdojo/dsge-svt
- Readily applied. Current application in progress, an open economy DSGE model that is about twice as big as the largest DSGE model estimated to date

References

- Chib & Ramamurthy (2010), "Tailored Randomized Block MCMC Methods with Application to DSGE Models," *Journal of Econometrics*
- Chib, Shin, & Tan (2021), "DSGE-SVt: An Econometric Toolkit for High-Dimensional DSGE Models with SV and t Errors," *Computational Economics*
- Sims (2001), "Solving Linear Rational Expectations Models," *Computational Economics*